

Richer Contexts

Maury Courtland (PhD; www.maury.science)

November 12, 2019

Talk Overview

Talk Overview

- ▶ 30k foot background

Talk Overview

- ▶ 30k foot background
- ▶ ELMo and deep representations

Talk Overview

- ▶ 30k foot background
- ▶ ELMo and deep representations
- ▶ BERT and pre-trained models

30k foot background: The benefits of pre-training and more information

30k foot background: The benefits of pre-training and more information

- ▶ Pre-training avoids repeated, costly spin-up time (and core-hours)

30k foot background: The benefits of pre-training and more information

- ▶ Pre-training avoids repeated, costly spin-up time (and core-hours)
 - ▶ Arguably what allowed CV to take off (AlexNet/ImageNet)

30k foot background: The benefits of pre-training and more information

- ▶ Pre-training avoids repeated, costly spin-up time (and core-hours)
 - ▶ Arguably what allowed CV to take off (AlexNet/ImageNet)
- ▶ Fixed models provide a common out-of-the-box starter model to extend (like a base-class module)

30k foot background: The benefits of pre-training and more information

- ▶ Pre-training avoids repeated, costly spin-up time (and core-hours)
 - ▶ Arguably what allowed CV to take off (AlexNet/ImageNet)
- ▶ Fixed models provide a common out-of-the-box starter model to extend (like a base-class module)
 - ▶ Helps researchers have a shared common ground of tools

30k foot background: The benefits of pre-training and more information

- ▶ Pre-training avoids repeated, costly spin-up time (and core-hours)
 - ▶ Arguably what allowed CV to take off (AlexNet/ImageNet)
- ▶ Fixed models provide a common out-of-the-box starter model to extend (like a base-class module)
 - ▶ Helps researchers have a shared common ground of tools
- ▶ What could be wrong with *more* information (**a little troll-y**)

30k foot background: The benefits of pre-training and more information

- ▶ Pre-training avoids repeated, costly spin-up time (and core-hours)
 - ▶ Arguably what allowed CV to take off (AlexNet/ImageNet)
- ▶ Fixed models provide a common out-of-the-box starter model to extend (like a base-class module)
 - ▶ Helps researchers have a shared common ground of tools
- ▶ What could be wrong with *more* information (**a little troll-y**)
 - ▶ More signal (as long as it's worth the extra learning time and complexity) can immensely boost performance

Peters et al. 2018: Deep contextualized word representations (aka ELMo)

Background and Approach

- ▶ Wants to capture both syntax and semantics as previously done

Background and Approach

- ▶ Wants to capture both syntax and semantics as previously done
- ▶ Wants to capture variance across contexts (i.e. polysemy) *not* previously done

Background and Approach

- ▶ Wants to capture both syntax and semantics as previously done
- ▶ Wants to capture variance across contexts (i.e. polysemy) *not* previously done
 - ▶ Tokens are assigned representations based on entire sentence

Background and Approach

- ▶ Wants to capture both syntax and semantics as previously done
- ▶ Wants to capture variance across contexts (i.e. polysemy) *not* previously done
 - ▶ Tokens are assigned representations based on entire sentence
 - ▶ Uses bidirectional LSTM with language modeling (LM) objective

Background and Approach

- ▶ Wants to capture both syntax and semantics as previously done
- ▶ Wants to capture variance across contexts (i.e. polysemy) *not* previously done
 - ▶ Tokens are assigned representations based on entire sentence
 - ▶ Uses bidirectional LSTM with language modeling (LM) objective
 - ▶ Thus, ELMo (**E**MBEDDINGS from **L**ANGUAGE **M**ODELS)

Strengths

- ▶ ELMo representations are deep (a function of all biLM layers, not just top layer)

Strengths

- ▶ ELMo representations are deep (a function of all biLM layers, not just top layer)
- ▶ Allows distributed meaning encoding

Strengths

- ▶ ELMo representations are deep (a function of all biLM layers, not just top layer)
- ▶ Allows distributed meaning encoding
 - ▶ Higher states capture context-dependent meaning (good for WSD)

Strengths

- ▶ ELMo representations are deep (a function of all biLM layers, not just top layer)
- ▶ Allows distributed meaning encoding
 - ▶ Higher states capture context-dependent meaning (good for WSD)
 - ▶ Lower states capture aspects of syntax (good for POS tagging)

Strengths

- ▶ ELMo representations are deep (a function of all biLM layers, not just top layer)
- ▶ Allows distributed meaning encoding
 - ▶ Higher states capture context-dependent meaning (good for WSD)
 - ▶ Lower states capture aspects of syntax (good for POS tagging)
 - ▶ With access to all this information, downstream models can select the relevant dimensions of information for the task (i.e. they are transferable)

Method

- ▶ Start with forward and backward LMs

Method

- ▶ Start with forward and backward LMs
 - ▶ Forward model: $p(t_1, t_2, \dots, t_N) = \prod_{k=1}^N p(t_k | t_1, t_2, \dots, t_{k-1})$

Method

- ▶ Start with forward and backward LMs
 - ▶ Forward model: $p(t_1, t_2, \dots, t_N) = \prod_{k=1}^N p(t_k | t_1, t_2, \dots, t_{k-1})$
 - ▶ Backward model:
 $p(t_1, t_2, \dots, t_N) = \prod_{k=1}^N p(t_k | t_{k+1}, t_{k+2}, \dots, t_N)$

Method

- ▶ Start with forward and backward LMs
 - ▶ Forward model: $p(t_1, t_2, \dots, t_N) = \prod_{k=1}^N p(t_k | t_1, t_2, \dots, t_{k-1})$
 - ▶ Backward model:
$$p(t_1, t_2, \dots, t_N) = \prod_{k=1}^N p(t_k | t_{k+1}, t_{k+2}, \dots, t_N)$$
- ▶ Jointly maximize their log likelihoods:

Method

- ▶ Start with forward and backward LMs
 - ▶ Forward model: $p(t_1, t_2, \dots, t_N) = \prod_{k=1}^N p(t_k | t_1, t_2, \dots, t_{k-1})$
 - ▶ Backward model:
 $p(t_1, t_2, \dots, t_N) = \prod_{k=1}^N p(t_k | t_{k+1}, t_{k+2}, \dots, t_N)$
- ▶ Jointly maximize their log likelihoods:

$$\sum_{k=1}^N (\log p(t_k \mid t_1, \dots, t_{k-1}; \Theta_x, \vec{\Theta}_{LSTM}, \Theta_s) \\ + \log p(t_k \mid t_{k+1}, \dots, t_N; \Theta_x, \overleftarrow{\Theta}_{LSTM}, \Theta_s))$$

ELMo's Magic



ELMo's Magic



...



Just combine all the representations computed for each token into a single vector!

each token t_k , a L -layer biLM computes a set of $2L + 1$ representations

$$\begin{aligned} R_k &= \{\mathbf{x}_k^{LM}, \overrightarrow{\mathbf{h}}_{k,j}^{LM}, \overleftarrow{\mathbf{h}}_{k,j}^{LM} \mid j = 1, \dots, L\} \\ &= \{\mathbf{h}_{k,j}^{LM} \mid j = 0, \dots, L\}, \end{aligned}$$

where $\mathbf{h}_{k,0}^{LM}$ is the token layer and $\mathbf{h}_{k,j}^{LM} = [\overrightarrow{\mathbf{h}}_{k,j}^{LM}; \overleftarrow{\mathbf{h}}_{k,j}^{LM}]$, for each biLSTM layer.

ELMo's Magic

- ▶ ...
- ▶ Just combine all the representations computed for each token into a single vector!

each token t_k , a L -layer biLM computes a set of $2L + 1$ representations

$$\begin{aligned} R_k &= \{ \mathbf{x}_k^{LM}, \overrightarrow{\mathbf{h}}_{k,j}^{LM}, \overleftarrow{\mathbf{h}}_{k,j}^{LM} \mid j = 1, \dots, L \} \\ &= \{ \mathbf{h}_{k,j}^{LM} \mid j = 0, \dots, L \}, \end{aligned}$$

where $\mathbf{h}_{k,0}^{LM}$ is the token layer and $\mathbf{h}_{k,j}^{LM} = [\overrightarrow{\mathbf{h}}_{k,j}^{LM}; \overleftarrow{\mathbf{h}}_{k,j}^{LM}]$, for each biLSTM layer.

- ▶ Previous approaches for contextually aware embeddings use a forward LM and take the last layer of the hidden state to be the embedding

Choosing a Useful Representation

- ▶ Simplest case is to simply select the top layer (as in previous work)

Choosing a Useful Representation

- ▶ Simplest case is to simply select the top layer (as in previous work)
- ▶ *Or*, for each task, adjust the weighting function of each layer:

$$\mathbf{ELMo}_k^{task} = E(R_k; \Theta^{task}) = \gamma^{task} \sum_{j=0}^L s_j^{task} \mathbf{h}_{k,j}^{LM}.$$

Choosing a Useful Representation

- ▶ Simplest case is to simply select the top layer (as in previous work)
- ▶ **Or**, for each task, adjust the weighting function of each layer:

$$\mathbf{ELMo}_k^{task} = E(R_k; \Theta^{task}) = \gamma^{task} \sum_{j=0}^L s_j^{task} \mathbf{h}_{k,j}^{LM}.$$

- ▶ where s^{task} are softmax-normalized weights and γ^{task} allows the downstream model to scale the ELMo vector.

Choosing a Useful Representation

- ▶ Simplest case is to simply select the top layer (as in previous work)
- ▶ *Or*, for each task, adjust the weighting function of each layer:

$$\mathbf{ELMo}_k^{task} = E(R_k; \Theta^{task}) = \gamma^{task} \sum_{j=0}^L s_j^{task} \mathbf{h}_{k,j}^{LM}.$$

- ▶ where s^{task} are softmax-normalized weights and γ^{task} allows the downstream model to scale the ELMo vector.
- ▶ sometimes it helps to apply layer normalization before weighting

Applying the Representations Downstream

- ▶ Use the pre-trained biLM and a specified architecture for target task

Applying the Representations Downstream

- ▶ Use the pre-trained biLM and a specified architecture for target task
 - 1) Run pre-trained (frozen) biLM and store representations for each word

Applying the Representations Downstream

- ▶ Use the pre-trained biLM and a specified architecture for target task
 - 1) Run pre-trained (frozen) biLM and store representations for each word
 - 2) Let target task model learn a linear combination of these representations

Applying the Representations Downstream

- ▶ Use the pre-trained biLM and a specified architecture for target task
 - 1) Run pre-trained (frozen) biLM and store representations for each word
 - 2) Let target task model learn a linear combination of these representations
- ▶ Can include (by concatenation) ELMo representations both in the input (concatenated with a context-independent embedding) and output (concatenated with the top hidden layer)

Applying the Representations Downstream

- ▶ Use the pre-trained biLM and a specified architecture for target task
 - 1) Run pre-trained (frozen) biLM and store representations for each word
 - 2) Let target task model learn a linear combination of these representations
- ▶ Can include (by concatenation) ELMo representations both in the input (concatenated with a context-independent embedding) and output (concatenated with the top hidden layer)
- ▶ Authors find it helpful to add “a moderate amount” of dropout

Applying the Representations Downstream

- ▶ Use the pre-trained biLM and a specified architecture for target task
 - 1) Run pre-trained (frozen) biLM and store representations for each word
 - 2) Let target task model learn a linear combination of these representations
- ▶ Can include (by concatenation) ELMo representations both in the input (concatenated with a context-independent embedding) and output (concatenated with the top hidden layer)
- ▶ Authors find it helpful to add “a moderate amount” of dropout
- ▶ Also some cases regularization helped (adding $\lambda ||w||_2^2$ to the loss making it close to an average of all layers)

Results: SOTA!

TASK	PREVIOUS SOTA		OUR BASELINE	ELMo + BASELINE	INCREASE (ABSOLUTE/ RELATIVE)
SQuAD	Liu et al. (2017)	84.4	81.1	85.8	4.7 / 24.9%
SNLI	Chen et al. (2017)	88.6	88.0	88.7 \pm 0.17	0.7 / 5.8%
SRL	He et al. (2017)	81.7	81.4	84.6	3.2 / 17.2%
Coref	Lee et al. (2017)	67.2	67.2	70.4	3.2 / 9.8%
NER	Peters et al. (2017)	91.93 \pm 0.19	90.15	92.22 \pm 0.10	2.06 / 21%
SST-5	McCann et al. (2017)	53.7	51.4	54.7 \pm 0.5	3.3 / 6.8%

Table 1: Test set comparison of ELMo enhanced neural models with state-of-the-art single model baselines across six benchmark NLP tasks. The performance metric varies across tasks – accuracy for SNLI and SST-5; F_1 for SQuAD, SRL and NER; average F_1 for Coref. Due to the small test sizes for NER and SST-5, we report the mean and standard deviation across five runs with different random seeds. The “increase” column lists both the absolute and relative improvements over our baseline.

Figure 1: ELMo’s Improvements Across the board

Analysis

- ▶ ELMo improves over just last layer and is better when linear weights are allowed to vary $\lambda = .001$ (better than $\lambda = 1$)




Figure 1: Comparison of baseline vs. ELMo performance for SNLI and SRL as the training set size is varied from 0.1% to 100%.

Figure 2: ELMo is very useful for small training corpora

Analysis

- ▶ ELMo improves over just last layer and is better when linear weights are allowed to vary $\lambda = .001$ (better than $\lambda = 1$)
- ▶ Isolation of layers as features for WSD and POS tagging tasks reveal different layers capture different information




Figure 1: Comparison of baseline vs. ELMo performance for SNLI and SRL as the training set size is varied from 0.1% to 100%.

Figure 2: ELMo is very useful for small training corpora

Analysis

- ▶ ELMo improves over just last layer and is better when linear weights are allowed to vary $\lambda = .001$ (better than $\lambda = 1$)
- ▶ Isolation of layers as features for WSD and POS tagging tasks reveal different layers capture different information
- ▶ ELMo also converges (at higher performance) much quicker than previous approaches




Figure 1: Comparison of baseline vs. ELMo performance for SNLI and SRL as the training set size is varied from 0.1% to 100%.

Figure 2: ELMo is very useful for small training corpora

Devlin et al. 2018: BERT: Pre-training of Deep
Bidirectional Transformers for Language Understanding

Background and Approach

- ▶ Pre-trained LMs improve performance on many NLP tasks (both sentence and token level)

Background and Approach

- ▶ Pre-trained LMs improve performance on many NLP tasks (both sentence and token level)
- ▶ 2 current approaches:

Background and Approach

- ▶ Pre-trained LMs improve performance on many NLP tasks (both sentence and token level)
- ▶ 2 current approaches:
 - 1) feature-based: use (frozen) pre-trained models as feature encoders and let downstream models learn the task, similar to embeddings

Background and Approach

- ▶ Pre-trained LMs improve performance on many NLP tasks (both sentence and token level)
- ▶ 2 current approaches:
 - 1) feature-based: use (frozen) pre-trained models as feature encoders and let downstream models learn the task, similar to embeddings
 - 2) fine-tuning: fine-tune pre-trained model parameters directly on downstream task

Background and Approach

- ▶ Pre-trained LMs improve performance on many NLP tasks (both sentence and token level)
- ▶ 2 current approaches:
 - 1) feature-based: use (frozen) pre-trained models as feature encoders and let downstream models learn the task, similar to embeddings
 - 2) fine-tuning: fine-tune pre-trained model parameters directly on downstream task
- ▶ same pre-train objective: “unidirectional” LMs, which limits architecture choices

Strengths

- ▶ Defines new training objective: “masked language model” (inspired by Cloze)

Strengths

- ▶ Defines new training objective: “masked language model” (inspired by Cloze)
 - ▶ Allows new architectures (here, deep bidirectional Transformer) given freedom from unidirectionality LM constraint (**BIG DEPARTURE** from past approaches)

Strengths

- ▶ Defines new training objective: “masked language model” (inspired by Cloze)
 - ▶ Allows new architectures (here, deep bidirectional Transformer) given freedom from unidirectionality LM constraint (**BIG DEPARTURE** from past approaches)
 - ▶ Randomly masks some input tokens (15%), task is to predict vocab ID of those tokens

Strengths

- ▶ Defines new training objective: “masked language model” (inspired by Cloze)
 - ▶ Allows new architectures (here, deep bidirectional Transformer) given freedom from unidirectionality LM constraint (**BIG DEPARTURE** from past approaches)
 - ▶ Randomly masks some input tokens (15%), task is to predict vocab ID of those tokens
 - ▶ Ex. from original Taylor, 1953: “Chickens cackle and _____ quack”

Strengths

- ▶ Defines new training objective: “masked language model” (inspired by Cloze)
 - ▶ Allows new architectures (here, deep bidirectional Transformer) given freedom from unidirectionality LM constraint (**BIG DEPARTURE** from past approaches)
 - ▶ Randomly masks some input tokens (15%), task is to predict vocab ID of those tokens
 - ▶ Ex. from original Taylor, 1953: “Chickens cackle and _____ quack”
- ▶ Avoids training and engineering time for task-specific architectures

Design Philosophy

- ▶ Have minimal architectural changes from pre-trained to downstream task (minimal parameters learned from scratch)




Figure 3: Applying pre-trained BERT to downstream tasks

Methods

- Built out of transformers (from “Attention Is All You Need” (2017)):




Figure 4: Transformer architecture

Inputs

- Use WordPiece embeddings (essentially sub-word features) with 30k vocab, denoted E_{token}




Figure 5: BERT input representation

Inputs

- ▶ Use WordPiece embeddings (essentially sub-word features) with 30k vocab, denoted E_{token}
 - ▶ Add segment embedding (E_A or E_B) and position embeddings (E_i) to token embeddings




Figure 5: BERT input representation

Inputs

- ▶ Use WordPiece embeddings (essentially sub-word features) with 30k vocab, denoted E_{token}
 - ▶ Add segment embedding (E_A or E_B) and position embeddings (E_i) to token embeddings
- ▶ Use “aggregate sequence representation for classification tasks” token [CLS] (final hidden vector C)




Figure 5: BERT input representation

Inputs

- ▶ Use WordPiece embeddings (essentially sub-word features) with 30k vocab, denoted E_{token}
 - ▶ Add segment embedding (E_A or E_B) and position embeddings (E_i) to token embeddings
- ▶ Use “aggregate sequence representation for classification tasks” token [CLS] (final hidden vector C)
- ▶ Can handle either a single sentence or a pair (e.g. $\langle Q, A \rangle$), separated by [SEP] token (final hidden vector $T_{[SEP]}$)




Figure 5: BERT input representation

Outputs (training objective): Masked LM

- ▶ Final hidden vectors for masked tokens at final transformer layer

Outputs (training objective): Masked LM

- ▶ Final hidden vectors for masked tokens at final transformer layer
- ▶ Fed into output softmax over vocab (like normal) and minimize cross-entropy loss

Outputs (training objective): Masked LM

- ▶ Final hidden vectors for masked tokens at final transformer layer
- ▶ Fed into output softmax over vocab (like normal) and minimize cross-entropy loss
- ▶ *BUT*, this creates mismatch between pre-training and fine-tuning (directly against design philosophy)

Outputs (training objective): Masked LM

- ▶ Final hidden vectors for masked tokens at final transformer layer
- ▶ Fed into output softmax over vocab (like normal) and minimize cross-entropy loss
- ▶ *BUT*, this creates mismatch between pre-training and fine-tuning (directly against design philosophy)
 - ▶ Solution: don't always [MASK] tokens

Outputs (training objective): Masked LM

- ▶ Final hidden vectors for masked tokens at final transformer layer
- ▶ Fed into output softmax over vocab (like normal) and minimize cross-entropy loss
- ▶ *BUT*, this creates mismatch between pre-training and fine-tuning (directly against design philosophy)
 - ▶ Solution: don't always [MASK] tokens
 - ▶ To remedy, they use 80% [MASK], 10% random token, 10% unmasked correct token

Outputs: Next Sentence Prediction

- ▶ Masked LM is good for token-level tasks, but downstream fine-tuning tasks have relationships between sentences (e.g. question answering, natural language inference, etc.)

Outputs: Next Sentence Prediction

- ▶ Masked LM is good for token-level tasks, but downstream fine-tuning tasks have relationships between sentences (e.g. question answering, natural language inference, etc.)
- ▶ Therefore we need a way to evaluate and embed sentence relationships

Outputs: Next Sentence Prediction

- ▶ Masked LM is good for token-level tasks, but downstream fine-tuning tasks have relationships between sentences (e.g. question answering, natural language inference, etc.)
- ▶ Therefore we need a way to evaluate and embed sentence relationships
- ▶ Create an output hidden vector from a special appended [CLS] token and (binary) classify whether the 2 training sentences are related or not (final accuracy is 97+% on NSP)

Outputs: Next Sentence Prediction

- ▶ Masked LM is good for token-level tasks, but downstream fine-tuning tasks have relationships between sentences (e.g. question answering, natural language inference, etc.)
- ▶ Therefore we need a way to evaluate and embed sentence relationships
- ▶ Create an output hidden vector from a special appended [CLS] token and (binary) classify whether the 2 training sentences are related or not (final accuracy is 97+% on NSP)
 - ▶ 50% of the time, sample adjacent sentences from the corpus (i.e. related)

Outputs: Next Sentence Prediction

- ▶ Masked LM is good for token-level tasks, but downstream fine-tuning tasks have relationships between sentences (e.g. question answering, natural language inference, etc.)
- ▶ Therefore we need a way to evaluate and embed sentence relationships
- ▶ Create an output hidden vector from a special appended [CLS] token and (binary) classify whether the 2 training sentences are related or not (final accuracy is 97+% on NSP)
 - ▶ 50% of the time, sample adjacent sentences from the corpus (i.e. related)
 - ▶ 50% of the time, sample random sentences from the corpus (i.e. unrelated)

Outputs: Next Sentence Prediction

- ▶ Masked LM is good for token-level tasks, but downstream fine-tuning tasks have relationships between sentences (e.g. question answering, natural language inference, etc.)
- ▶ Therefore we need a way to evaluate and embed sentence relationships
- ▶ Create an output hidden vector from a special appended [CLS] token and (binary) classify whether the 2 training sentences are related or not (final accuracy is 97+% on NSP)
 - ▶ 50% of the time, sample adjacent sentences from the corpus (i.e. related)
 - ▶ 50% of the time, sample random sentences from the corpus (i.e. unrelated)
- ▶ Simple fix, but boosts performance across the board (esp. for QNLI: 3.5% and SQuAD: .6%)

Fine-tuning

- ▶ Attention mechanisms make BERT easily extensible and can capture many tasks (rather than learning representations from scratch, just change weighting of subspace features)

Fine-tuning

- ▶ Attention mechanisms make BERT easily extensible and can capture many tasks (rather than learning representations from scratch, just change weighting of subspace features)
- ▶ Just plug-n-play: use whichever task-specific inputs and outputs you want and fine-tune **all** the parameters

Fine-tuning

- ▶ Attention mechanisms make BERT easily extensible and can capture many tasks (rather than learning representations from scratch, just change weighting of subspace features)
- ▶ Just plug-n-play: use whichever task-specific inputs and outputs you want and fine-tune **all** the parameters
 - ▶ Input: Use A/B sentence pairs (e.g. sentence and paraphrase, hypothesis and premise, etc.) OR just a dummy null B sentence (e.g. text classification, sequence tagging)

Fine-tuning

- ▶ Attention mechanisms make BERT easily extensible and can capture many tasks (rather than learning representations from scratch, just change weighting of subspace features)
- ▶ Just plug-n-play: use whichever task-specific inputs and outputs you want and fine-tune **all** the parameters
 - ▶ Input: Use A/B sentence pairs (e.g. sentence and paraphrase, hypothesis and premise, etc.) OR just a dummy null B sentence (e.g. text classification, sequence tagging)
 - ▶ Output for token-tasks: use token hidden vectors as input for token-level tasks (e.g. sequence tagging, QA)

Fine-tuning

- ▶ Attention mechanisms make BERT easily extensible and can capture many tasks (rather than learning representations from scratch, just change weighting of subspace features)
- ▶ Just plug-n-play: use whichever task-specific inputs and outputs you want and fine-tune **all** the parameters
 - ▶ Input: Use A/B sentence pairs (e.g. sentence and paraphrase, hypothesis and premise, etc.) OR just a dummy null B sentence (e.g. text classification, sequence tagging)
 - ▶ Output for token-tasks: use token hidden vectors as input for token-level tasks (e.g. sequence tagging, QA)
 - ▶ Output for sentence-tasks: Use C for sentence-level tasks (e.g. entailment, sentiment analysis)

Fine-tuning

- ▶ Attention mechanisms make BERT easily extensible and can capture many tasks (rather than learning representations from scratch, just change weighting of subspace features)
- ▶ Just plug-n-play: use whichever task-specific inputs and outputs you want and fine-tune **all** the parameters
 - ▶ Input: Use A/B sentence pairs (e.g. sentence and paraphrase, hypothesis and premise, etc.) OR just a dummy null B sentence (e.g. text classification, sequence tagging)
 - ▶ Output for token-tasks: use token hidden vectors as input for token-level tasks (e.g. sequence tagging, QA)
 - ▶ Output for sentence-tasks: Use C for sentence-level tasks (e.g. entailment, sentiment analysis)
- ▶ Relatively inexpensive (all paper tasks fine-tuned in a few hours on 1 GPU)

Results: Grand-slam SOTA

System	MNLI-(m/mm)	QQP	QNLI	SST-2	CoLA	STS-B	MRPC	RTE	Average
	392k	363k	108k	67k	8.5k	5.7k	3.5k	2.5k	-
Pre-OpenAI SOTA	80.6/80.1	66.1	82.3	93.2	35.0	81.0	86.0	61.7	74.0
BiLSTM+ELMo+Attn	76.4/76.1	64.8	79.8	90.4	36.0	73.3	84.9	56.8	71.0
OpenAI GPT	82.1/81.4	70.3	87.4	91.3	45.4	80.0	82.3	56.0	75.1
BERT _{BASE}	84.6/83.4	71.2	90.5	93.5	52.1	85.8	88.9	66.4	79.6
BERT _{LARGE}	86.7/85.9	72.1	92.7	94.9	60.5	86.5	89.3	70.1	82.1

Figure 6: SOTA across the board with some impressive gains

Ablation studies takeaways

- ▶ NSP fairly important (especially for QA and NLI)

Ablation studies takeaways

- ▶ NSP fairly important (especially for QA and NLI)
- ▶ Deep bidirectionality very important

Ablation studies takeaways

- ▶ NSP fairly important (especially for QA and NLI)
- ▶ Deep bidirectionality very important
 - ▶ Also cheaper than separate unidirectional models

Ablation studies takeaways

- ▶ NSP fairly important (especially for QA and NLI)
- ▶ Deep bidirectionality very important
 - ▶ Also cheaper than separate unidirectional models
 - ▶ And strictly more powerful than unidirectional models

Ablation studies takeaways

- ▶ NSP fairly important (especially for QA and NLI)
- ▶ Deep bidirectionality very important
 - ▶ Also cheaper than separate unidirectional models
 - ▶ And strictly more powerful than unidirectional models
- ▶ More parameters = better (sometimes *much* = 8%, sometimes *a bit* .4%)

Ablation studies takeaways

- ▶ NSP fairly important (especially for QA and NLI)
- ▶ Deep bidirectionality very important
 - ▶ Also cheaper than separate unidirectional models
 - ▶ And strictly more powerful than unidirectional models
- ▶ More parameters = better (sometimes *much* = 8%, sometimes *a bit* .4%)
- ▶ ELMo-ish approach of concatenating parameters as features is worse (though not by much, best linear combo is .3 worse on 96.1 F1) than a fine-tuning approach (which again, is generally cheaper and more portable)